

# Model-Based RL

---

Reinforcement Learning  
School of Data Science  
University of Virginia

Last updated: July 21, 2025

# Agenda

- > Model-Free RL
- > Model-Based RL
- > World Models

# Model Free vs. Model-Based

So far we have looked at model-free approaches

There was no transition model  $P(s_{t+1} | s_t, a_t)$

Instead, we *sampled* next state by running action in environment:

$$(s_t, a_t) \xrightarrow{\text{environment}} s_{t+1}, r_{t+1}$$

# Reminder about Model-Free RL

In some cases, we devised toy rules

In other cases, we ran a simulator



# Reminder about Model-Free RL

In some cases, we devised toy rules

In other cases, we ran a simulator

Of course, a simulator might use a model

But the RL agent doesn't know or learn the model



# Disadvantages of Model-Free RL

> **Lower sample efficiency:** without a model, agent only learns from experience

# Disadvantages of Model-Free RL

- > **Lower sample efficiency:** without a model, agent only learns from experience
- > **Can't plan ahead:** it is not possible to simulate rollouts and learn from them



# Disadvantages of Model-Free RL

- > **Lower sample efficiency:** without a model, agent only learns from experience
- > **Can't plan ahead:** it is not possible to simulate rollouts and learn from them
- > **Adaptability is challenging:** if environment / reward function changes, a lot of experience is required for learning



# Disadvantages of Model-Free RL

- > **Lower sample efficiency:** without a model, agent only learns from experience
- > **Can't plan ahead:** it is not possible to simulate rollouts and learn from them
- > **Adaptability is challenging:** if environment / reward function changes, a lot of experience is required for learning
- > **Lack of interpretability:** value functions and policies can be black boxes

# Model-Based RL – General Strategy

Two parts:

1. Learn a **dynamics function** to model observed state transitions  $P(s_{t+1} | s_t, a_t)$
2. Use model predictions to learn what **actions** to take (e.g., learn a policy)

# Action Selection and Horizon


We can use model predictions to learn what **actions** to take

For example, our usual strategy is to maximize return (= expected total discounted reward)

**Q: What horizon to use?**

Infinite horizon won't work ... prediction errors can compound

Instead, we use some planning horizon  $H$



# Action Selection with Finite Planning Horizon

1. Set planning horizon  $H$
2. Generate  $K$  random action sequences each with length  $H$ , denoted

$$\mathbf{A}^{(k)} = (a_t^{(k)}, \dots, a_{t+H-1}^{(k)})$$

3. Use the dynamics model  $f_\theta$  to predict the future states after taking each action sequence
4. Evaluate the return associated with each candidate action sequence
5. Select the best action sequence

This method is called *random shooting*

# Refinement to Action Selection: Replanning

Since our model is imperfect, we might have compounding errors as we plan into the future

We can adopt a *model predictive control* (MPC) approach:

1. At each time step, we perform random shooting or something else
2. Select the best  $H$ -step action sequence
3. Only take the first action from the sequence
4. Now replan at the next time step using updated state information

# Refinement to Dynamics Model: Ensemble

We are using some model (e.g., a neural network)  $f_{\theta}$  to predict the next state

A method for **potentially improving predictions** is to use a set of models  $\{f_{\theta_n}\}_{n=1}^N$

These can be independently initialized

**At inference time:** For each candidate action sequence, generate  $N$  independent rollouts

Average the rewards of the rollouts to select the best sequence

# Model-Based RL

Q: Given model-free RL limitations, why don't we always use a model?

# Model-Based RL

Q: Given model-free RL limitations, why don't we always use a model?

**A: Because developing an accurate model can be hard.**

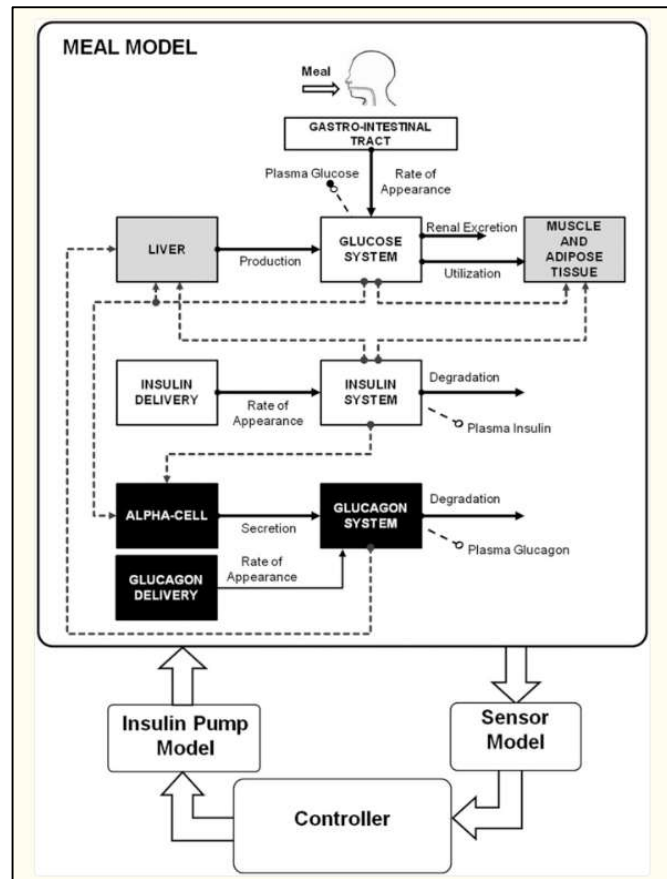
$$\begin{aligned}x_1[k + 1] = & \theta_1 u[k - 1] + \theta_2 u[k] + \theta_3 u[k + 1] \\& + \theta_4 x_1[k - 1] + \theta_5 x_1[k] + \theta_6 x_2[k] + \theta_0\end{aligned}$$

Individualization of pharmacological anemia management  
using reinforcement learning<sup>☆</sup>

Adam E. Gaweda<sup>a,\*</sup>, Mehmet K. Muezzinoglu<sup>b</sup>, George R. Aronoff<sup>a</sup>, Alfred A. Jacobs<sup>a</sup>,  
Jacek M. Zurada<sup>b</sup>, Michael E. Brier<sup>a,c</sup>



# Model-Based RL: Diabetes Simulator



► J Diabetes Sci Technol. 2014 Jan;8(1):26–34. doi: [10.1177/1932296813514502](https://doi.org/10.1177/1932296813514502)

## The UVA/PADOVA Type 1 Diabetes Simulator

New Features

[Chiara Dalla Man](#)<sup>1</sup>, [Francesco Micheletto](#)<sup>1</sup>, [Dayu Lv](#)<sup>2</sup>, [Marc Breton](#)<sup>2</sup>, [Boris Kovatchev](#)<sup>2</sup>, [Claudio Cobelli](#)<sup>1</sup>,

The model of glucose kinetics is described by,

$$\begin{cases} \dot{G}_p = EGP - U_{ii} - k_1 \cdot G_p(t) + k_2 \cdot G_t(t) & G_p(0) = G_{pb} \\ \dot{G}_t = -U_{id}(t) + k_1 \cdot G_p(t) - k_2 \cdot G_t(t) & G_t(0) = G_{pb} \frac{k_1}{k_2} \end{cases}$$

# World Models

# World Models

A *world model* (David Ha, Jürgen Schmidhuber, 2018) is a learned model of the dynamics:

Useful for:

$$(s_t, a_t) \rightarrow s_{t+1}, r_{t+1}$$

- > Predicting next state

- > Planning

- > Imagining possible futures without taking the actions in real environment.

**Critical when it would be dangerous or costly to try the actions in real world.**


**[worldmodels.github.io](https://worldmodels.github.io)**

# World Models – Classic Model

From Ha & Schmidhuber paper, architecture has three parts:

<b>VAE (encoder)</b>	Compress high-dim obs (e.g., images) into a low-dim latent space
<b>MDN-RNN</b>	Learn to predict next latent state, given current latent state, action
<b>Controller</b>	“Small” policy network that decides actions in latent space

# Car Racing Experiment - Image Encoding




Actual observations from the environment.

What gets encoded into  $z_t$ .

# World Model with Images

David Ha, Jürgen Schmidhuber, 2018



# MDN-RNN

Predicts next latent state as density function  $p(z)$

Approximated as mixture of Gaussians

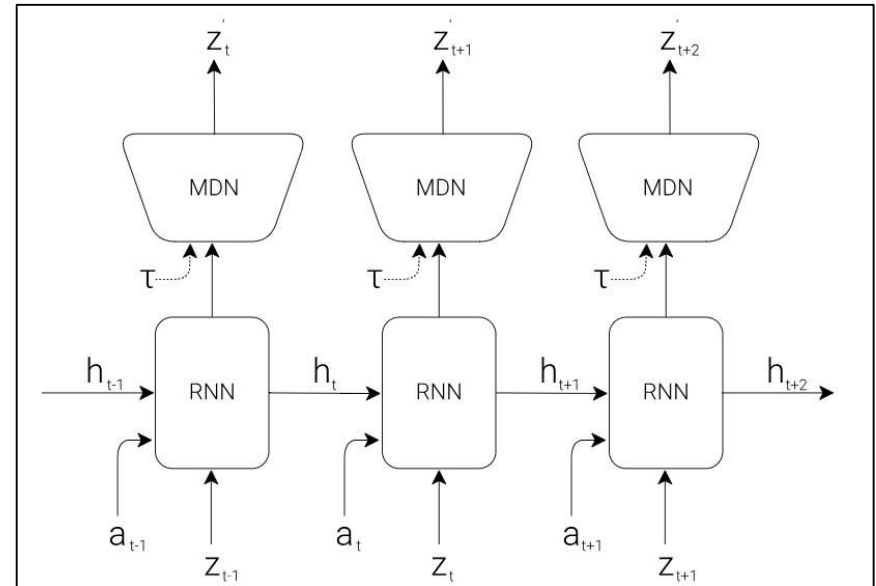
RNN used to model  $P(z_{t+1} \mid a_t, z_t, h_t)$

where  $a_t$  denotes the action

$h_t$  is the hidden state

$\tau$  is the *temperature* for controlling uncertainty

MDN component outputs parameters of mixture distr.



# Car Racing Experiment - Procedure

1. Collect 10,000 rollouts from a random policy.
2. Train VAE (V) to encode each frame into a latent vector  $z \in \mathcal{R}^{32}$ .
3. Train MDN-RNN (M) to model  $P(z_{t+1} \mid a_t, z_t, h_t)$ .
4. Evolve Controller (C) to maximize the expected cumulative reward of a rollout.



# Learning in a Dream

We have seen the procedure for training a simple policy to solve tasks

Can train the agent inside its “dream” environment

Then transfer policy back to actual environment

---

Explore the paper and interactive demo:

<https://worldmodels.github.io/>

